



Human trait grading benchmark

114 specimens; model run Qwen3-VL-32B-Instruct-FP8; prompt style sys4

Macro agreement



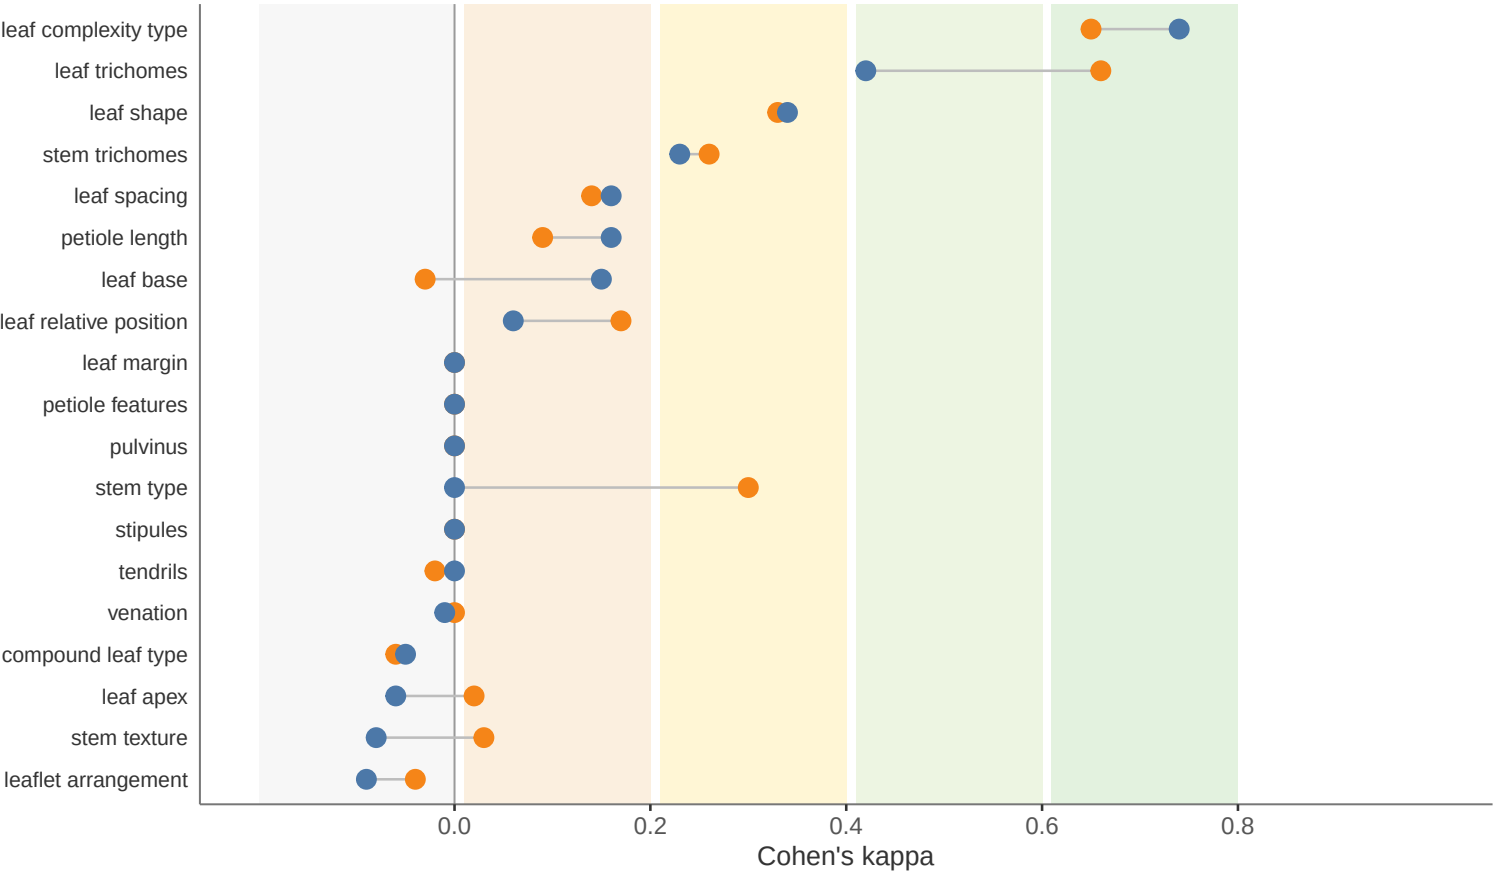
Macro kappa



Per-trait model agreement beyond chance

Cohen's kappa by trait; grey lines show the gap between model-Roni and model-Carmen agreement

Model vs Carmen Model vs Roni



Interpretation bands: slight 0–0.20, fair 0.21–0.40, moderate 0.41–0.60, substantial 0.61–0.80. Macro agreement/kappa are report-level averages across available pairwise scored items; the per-trait panel shows 19 model-scored traits from 21 extracted trait rows.