


SeedLearn trait grading: model agreement with expert and student references

Macro agreement is 68% for Model vs Roni and 68% for Model vs Carmen; human-human agreement is 79%. Traits shown: 19 model-scored traits, sorted by Model-vs-Roni kappa.



Source: human_grading_report.html, Qwen/Qwen3-VL-32B-Instruct-FP8 sys4 run, generated 2026-07-06. Extracted report has 21 trait rows; 19 include model-vs-human comparisons. Agreement is pairwise agreement, not ground-truth accuracy. Cohen's kappa is chance-corrected: 0 = chance-level, 0.4 ~ moderate, 0.8+ strong.